

Hadoop Core components and Daemons

2 components of Hadoop

1.HDFS(storage)

2.YARN/MRv2(processing)

Daemons

Namenode

Datanode

ResourceManger

NodeManger

Hadoop Core components and Daemons


2 components of Hadoop
Master-Slave

1.HDFS(storage)
Namenode(Master)
Datanode(Slave)

2.YARN(processing)
ResourceManger(Master)
NodeManager(Slave)

Simple Hadoop cluster with Daemons


Master




5 Servers in the cluster

1 – Master


4 – Slaves




Slave1



Slave2




Slave3





Slave n

Simple cluster with Hadoop Daemons

Master 1

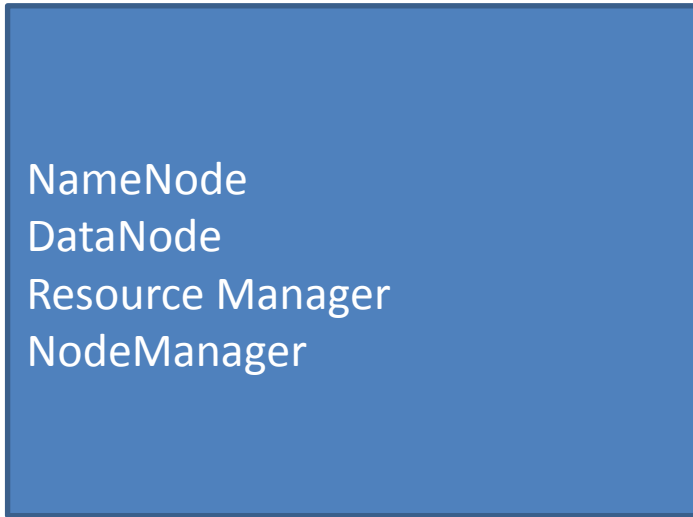


Master 2



Pseudo Distributed Mode

Single Server



File Blocks

By Default Block cutoff size is 128 MB

200mb – abc.txt

128mb – Block 1 

72mb -- Block 2 

500mb – weblog.dat

128mb – Block 1 

128mb -- Block 2 

128mb – Block 3 

116mb – Block 4 

Files with less than 128MB just takes one block of exact size.

100mb -- books.xml

Block1 

100 kb – flower.jpg

Block1 

File Blocks

First of all, if the block size were too small it would overwhelm the metadata server with too many blocks to track.

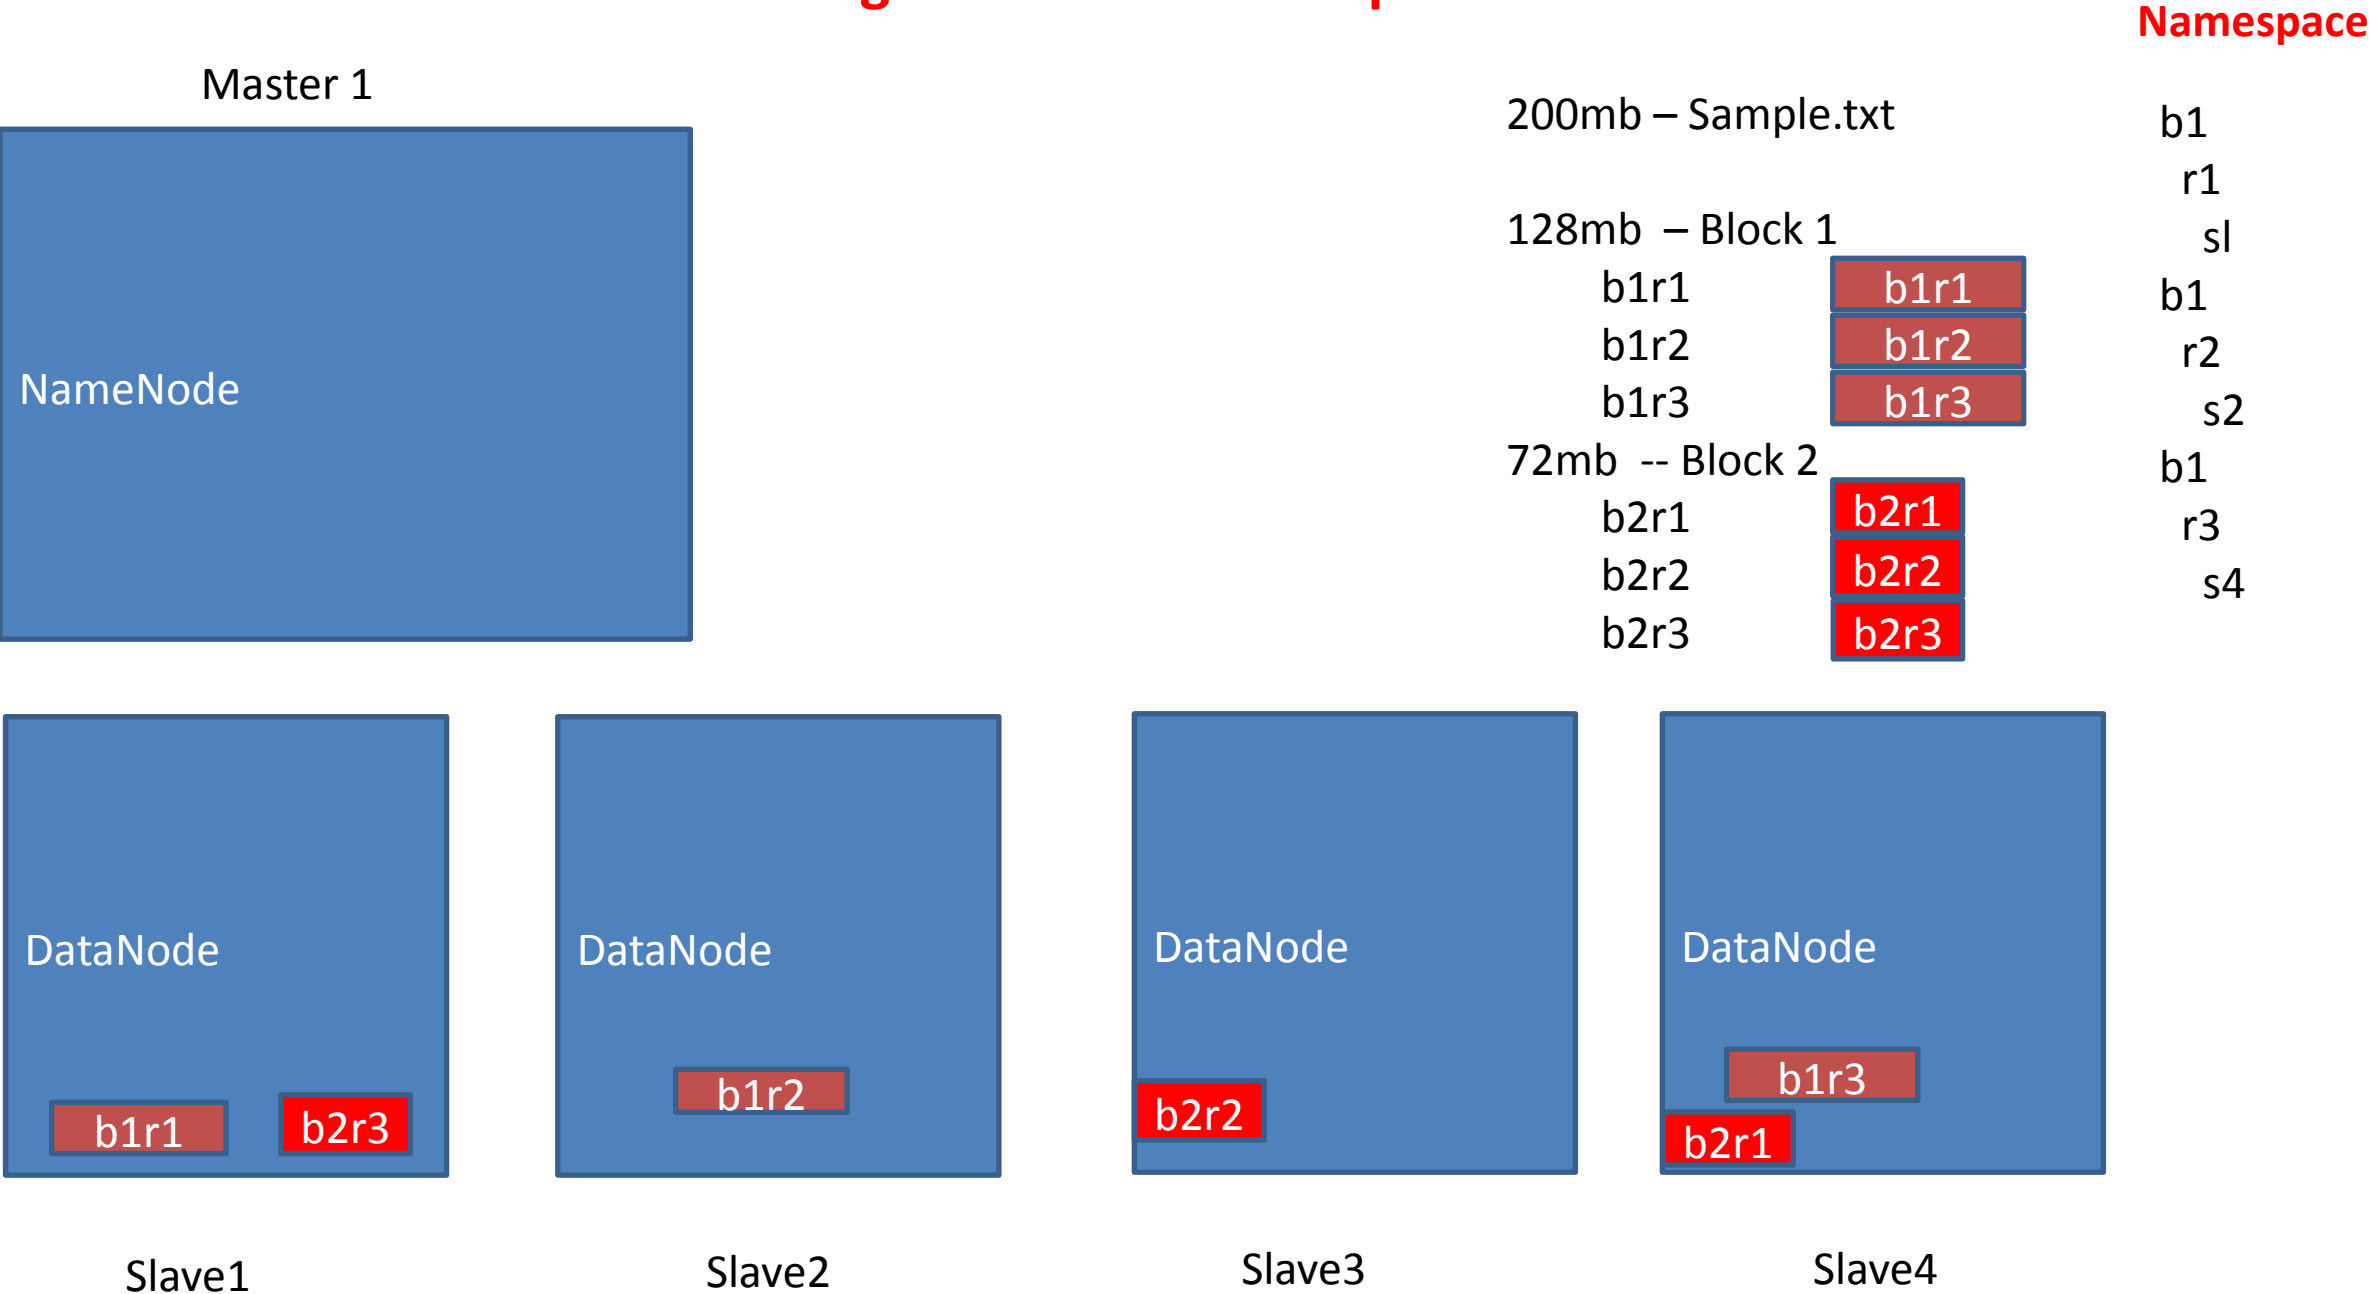
Second, HDFS is designed to enable high throughput so that the parallel processing of these large data sets happens as quickly as possible.

On one hand, the block size needs to be large enough to warrant the resources dedicated to an individual unit of data processing (for instance, a map or reduce task)

You want to find a balance where each task is able to process a reasonable amount of data while still getting the benefits of parallelism.

The smaller the block size, the more tasks you get, the more scheduling activity occurs.

Data Storage in Slaves with replication



MapReduce



Heartbeat: Confirm that the DataNode is operating and the block replicas it hosts are available

- **Total storage capacity**
- **Fraction of storage in use**
- **The number of data transfers currently in progress**

NameNode


These statistics are used for the NameNode's block allocation and load balancing decisions.

Heartbeat:
Default interval
is 3 seconds

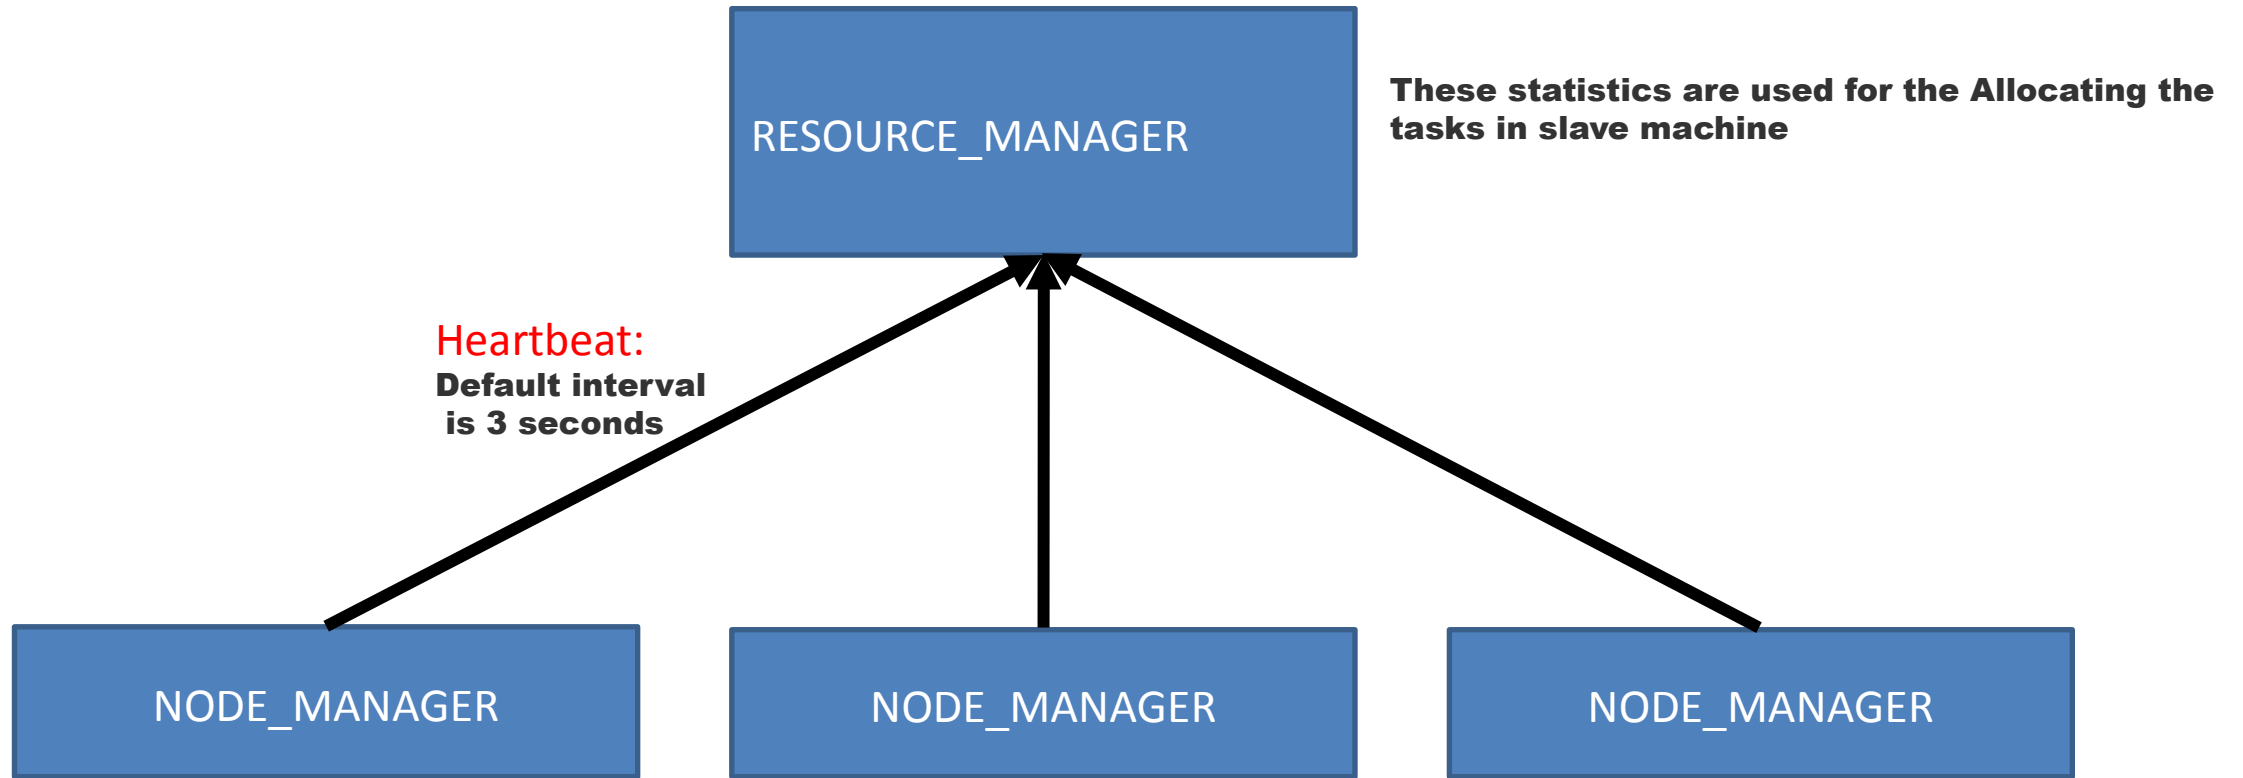
DataNode

DataNode


DataNode




Heartbeat: Confirm that the **NODEManager** is operating



UserMachine_1




UserMachine_2



Client Interface



NameNode
ResourceManager




Master

DataNode
NodeManager



Slave1

DataNode
NodeManager



Slave2

DataNode
NodeManager



Slave3

DataNode
NodeManager



Slave4

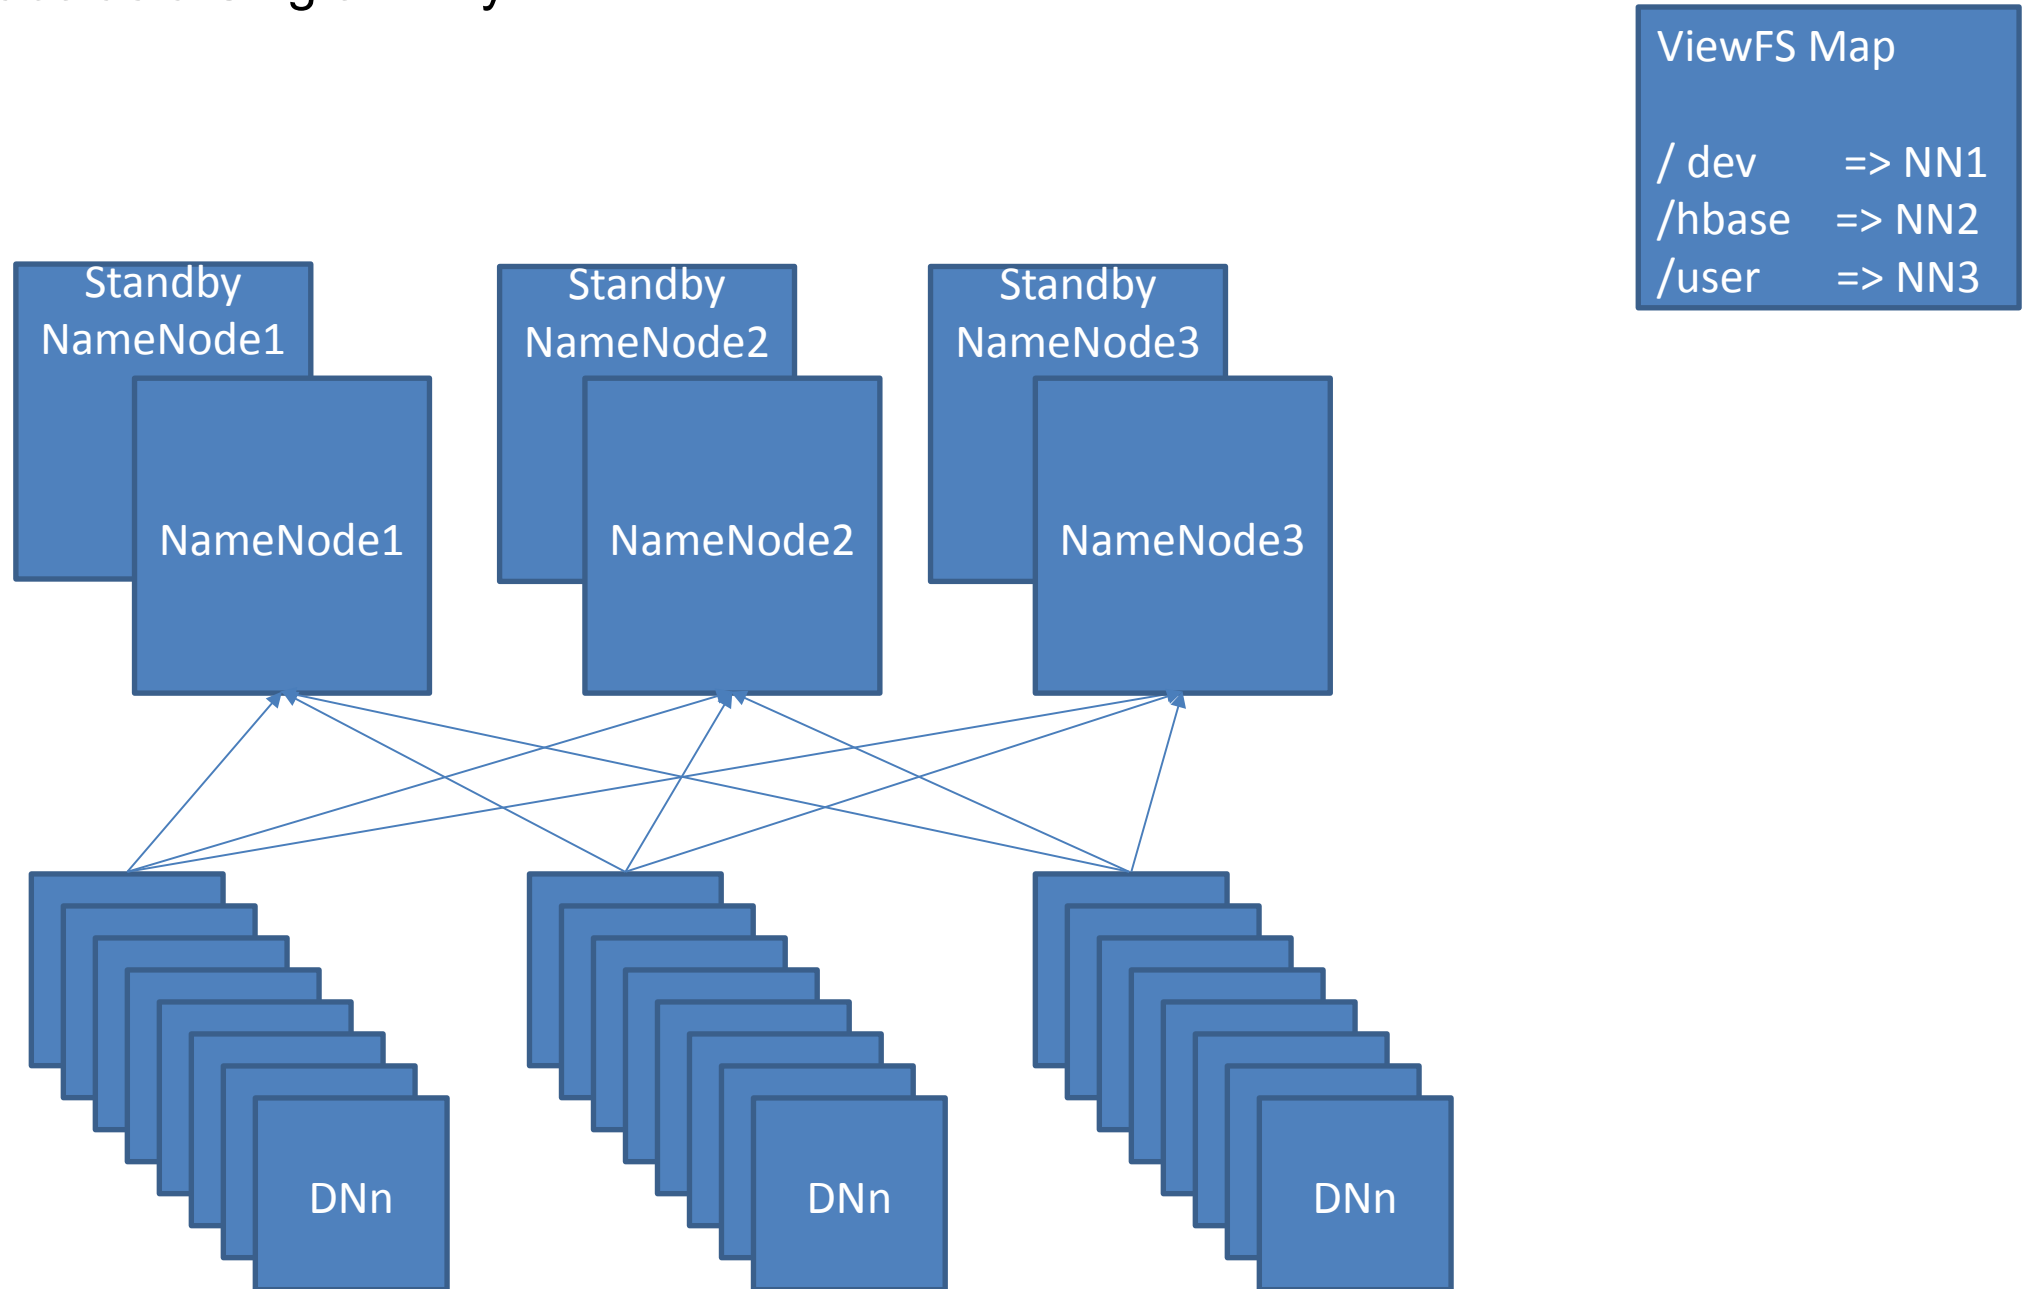
Load data in the HDFS cluster

Submit MapReduce jobs (describing how to process the data)

Retrieve or view the results of the job after its completion

Submit Pig or Hive queries

Client of HDFS uses a specialize plugin called viewFS to view the logical, global namespace as a Single Entity



```
<configuration>
```

```
<property>
```

```
<name>dfs.nameservices</name>
```

```
<value>dev,hbase,user</value>
```

```
</property>
```

```
<property>
```

```
<name>dfs.namenode.rpc-address.dev</name>
```

```
<value>nn-host1:rpc-port</value>
```

```
</property>
```

```
<property>
```

```
<name>dfs.namenode.rpc-address.hbase</name>
```

```
<value>nn-host2:rpc-port</value>
```

```
</property>
```

```
<property>
```

```
<name>dfs.namenode.rpc-address.user</name>
```

```
<value>nn-host3:rpc-port</value>
```

```
</property>
```

```
....
```


Other common configuration

Architecture of YARN

ResourceManager governs an entire cluster and manages the assignment of applications to underlying compute resources.

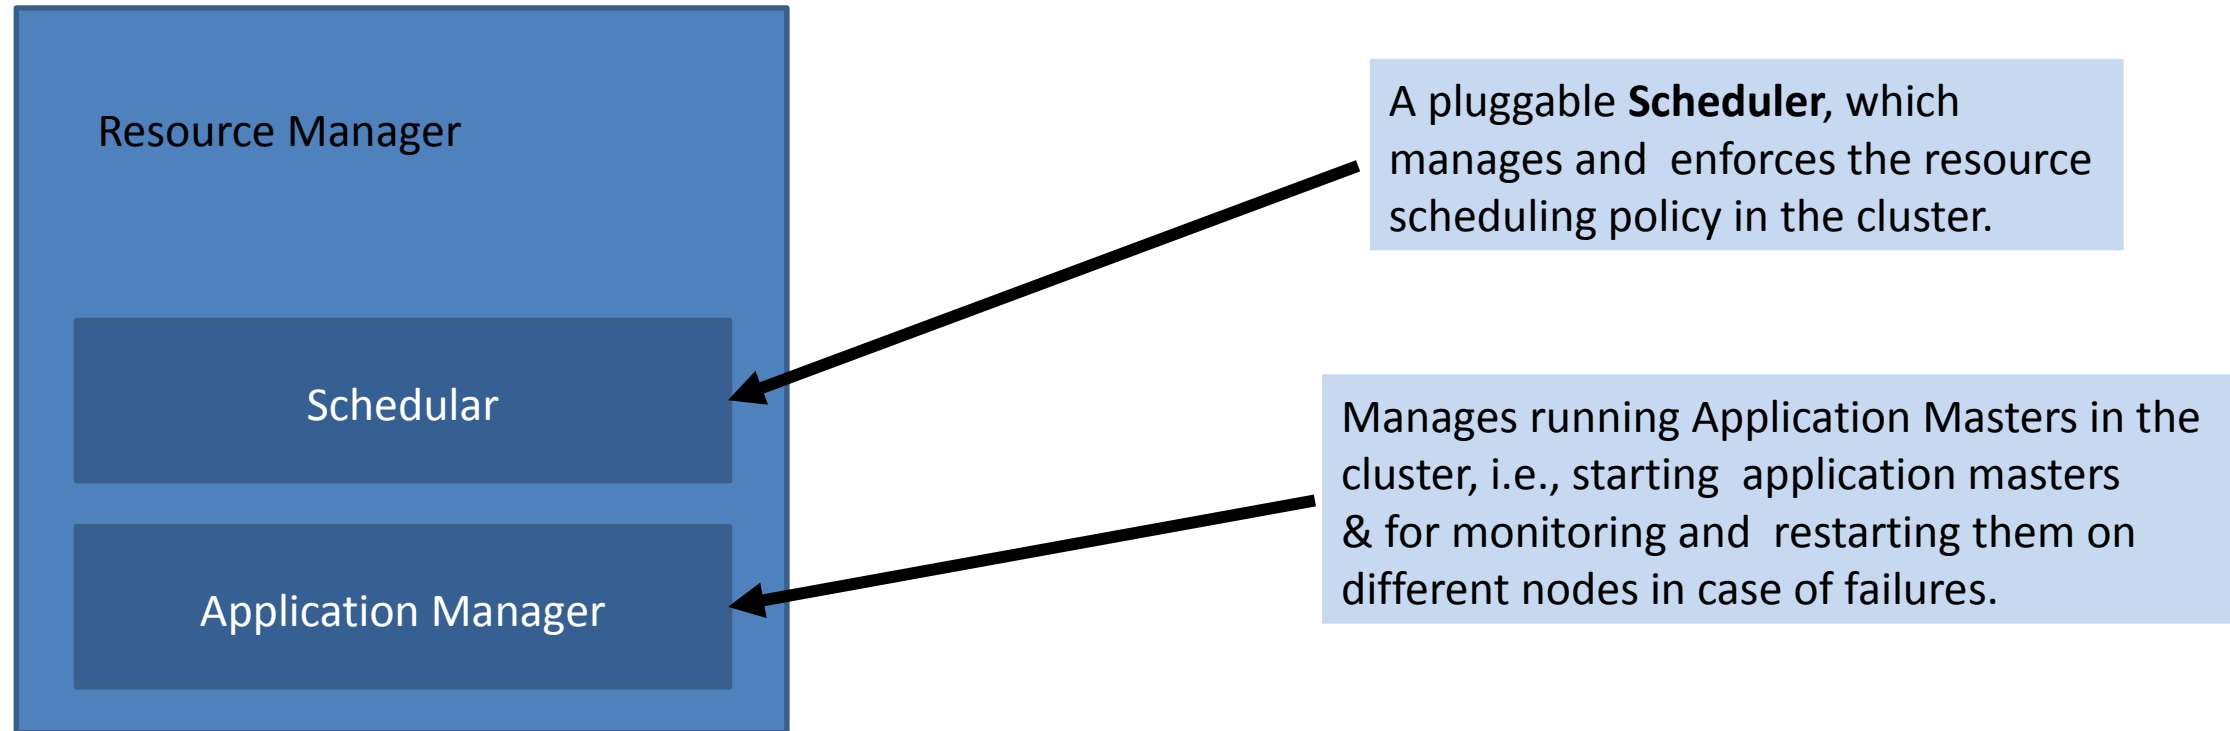
ApplicationMaster manages each instance of an application that runs within YARN.

NodeManager provides per-node services within the cluster, from overseeing the management of a container over its life cycle to monitoring resources and tracking the health of its node.



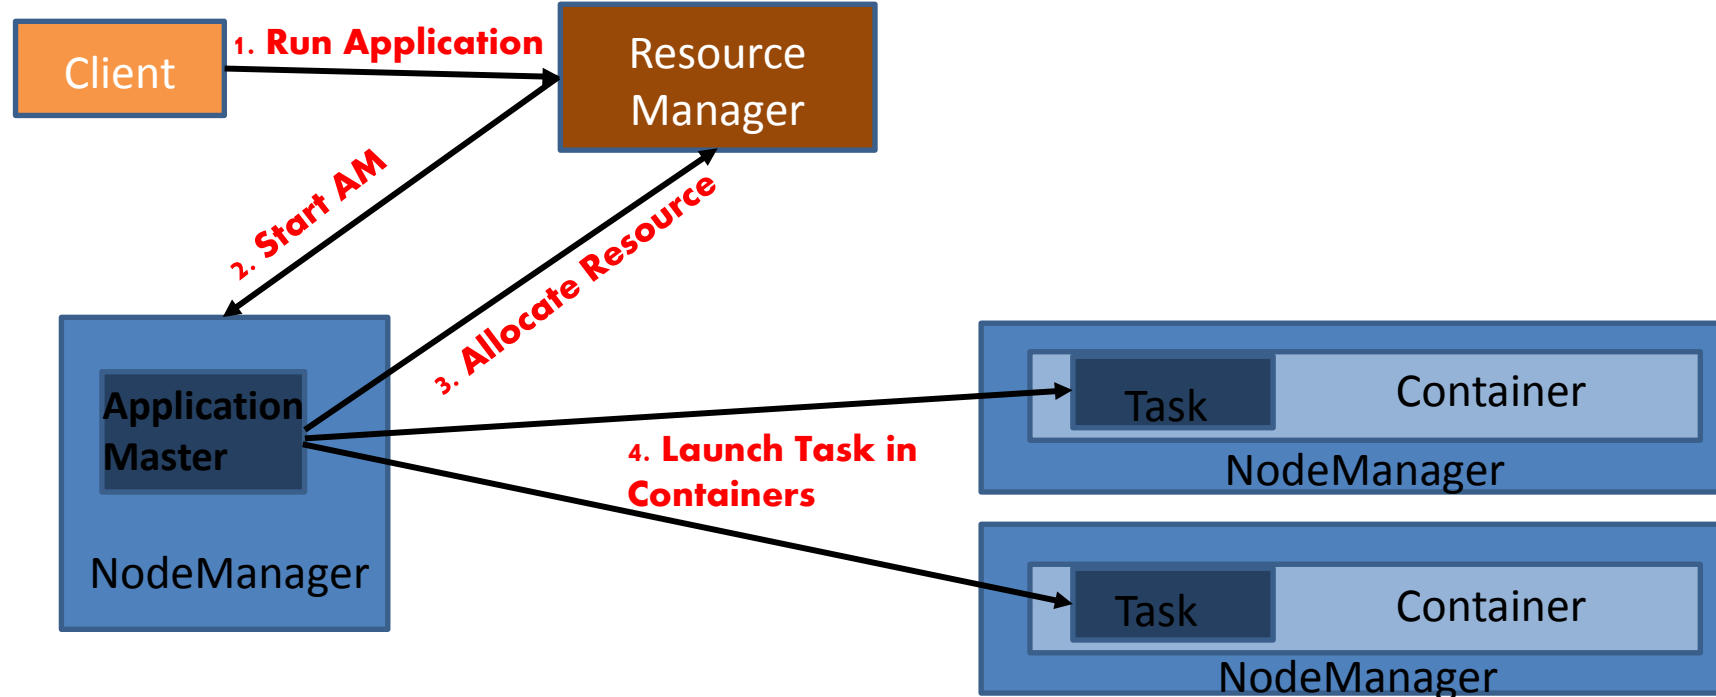
Resource Manager

There is a single Resource Manager, which has two main services:




Application Submission in YARN

1. Application Submission Client submits an Application to the YARN Resource Manager. The client needs to provide sufficient information to the ResourceManager in order to launch ApplicationMaster
2. YARN ResourceManager starts ApplicationMaster.
3. The ApplicationMaster then communicates with the Resource Manager to request resource allocation.
4. After a container is allocated to it, the ApplicationMaster communicates with the NodeManager to launch the tasks in the container.




Simple cluster with Hadoop Daemons

Master 1



Master 2



Master 3

